

PROBLEM STATEMENTS

This data was extracted from the 1994 Census bureau database by Ronny Kohavi and Barry Becker (Data Mining and Visualization, Silicon Graphics). A set of reasonably clean records was extracted using the following conditions: ((AAGE>16) & (AGI>100) & (AFNLWGT>1) && (HRSWK>0)). The prediction task is to determine whether a person makes over \$50K a year or less then.

Importing Libraries that will come handy for the projects

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

In [2]:

```
columns = ['Age', 'Work Class', 'Final Weight', 'Education', 'Education Number', 'Marital Status', 'Occupation', 'Relationship', 'Race', 'Sex', 'Capital Gain', 'Capital Loss', 'Hours per week', 'Income']
data = pd.read_csv("D:\\data set\\adult.data", names = columns)
```

In [3]:

```
data
```

Out[3]:

	Age	Work Class	Final Weight	Education	Education Number	Marital Status	Occupation	Relationship	Race	Sex	Income
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	Less than \$50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	\$50K and more
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	Less than \$50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	Less than \$50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	\$50K and more
...
32556	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	\$50K and more
32557	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	Less than \$50K

	Age	Work Class	Final Weight	Education	Education Number	Marital Status	Occupation	Relationship	Race	!
32558	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Ferr
32559	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	N
32560	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Ferr

32561 rows × 15 columns



From the above code we can observe that the data set contains 32561 rows × 15 columns.

Exploratory Data Analysis (EDA)

In [4]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   32561 non-null  int64
1   Work Class            32561 non-null  object
2   Final Weight          32561 non-null  int64
3   Education              32561 non-null  object
4   Education Number      32561 non-null  int64
5   Marital Status        32561 non-null  object
6   Occupation            32561 non-null  object
7   Relationship          32561 non-null  object
8   Race                  32561 non-null  object
9   Sex                   32561 non-null  object
10  Capital Gain          32561 non-null  int64
11  Capital Loss          32561 non-null  int64
12  Hours per Week        32561 non-null  int64
13  Country               32561 non-null  object
14  Income                32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

From the above we can observe that our data do not contain any missing value and our target value "Income" has an Object datatype, so before moving further lets convert it into numerical datatype

In [5]:

```
from sklearn.preprocessing import LabelEncoder

labelEncoder = LabelEncoder()
data['Income'] = labelEncoder.fit_transform(data['Income'])
```

In [6]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
```

Data columns (total 15 columns):

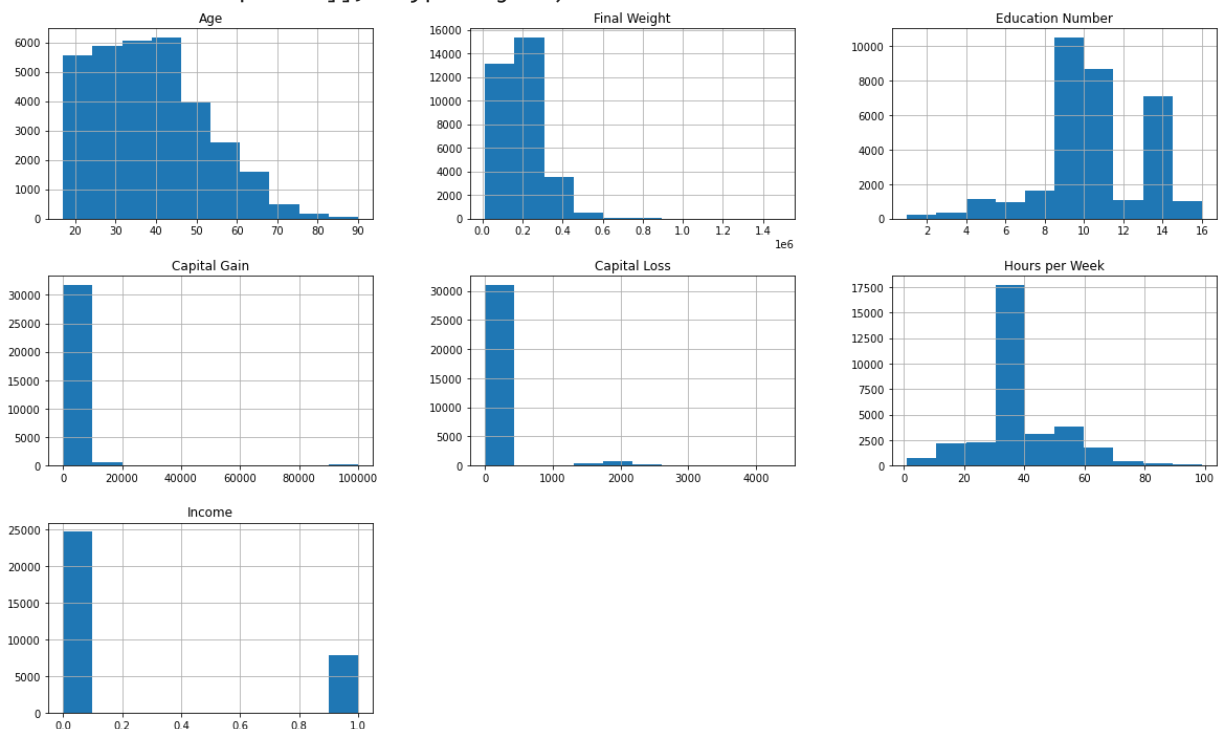
#	Column	Non-Null Count	Dtype
0	Age	32561 non-null	int64
1	Work Class	32561 non-null	object
2	Final Weight	32561 non-null	int64
3	Education	32561 non-null	object
4	Education Number	32561 non-null	int64
5	Marital Status	32561 non-null	object
6	Occupation	32561 non-null	object
7	Relationship	32561 non-null	object
8	Race	32561 non-null	object
9	Sex	32561 non-null	object
10	Capital Gain	32561 non-null	int64
11	Capital Loss	32561 non-null	int64
12	Hours per Week	32561 non-null	int64
13	Country	32561 non-null	object
14	Income	32561 non-null	int32

dtypes: int32(1), int64(6), object(8)

memory usage: 3.6+ MB

```
In [7]: data.hist(figsize=(20,12))
```

```
Out[7]: array([[<AxesSubplot:title={'center':'Age'}>,
      <AxesSubplot:title={'center':'Final Weight'}>,
      <AxesSubplot:title={'center':'Education Number'}>],
      [<AxesSubplot:title={'center':'Capital Gain'}>,
      <AxesSubplot:title={'center':'Capital Loss'}>,
      <AxesSubplot:title={'center':'Hours per Week'}>],
      [<AxesSubplot:title={'center':'Income'}>, <AxesSubplot:>,
      <AxesSubplot:>]], dtype=object)
```




From the above graphs we can conclude

a) Age and Hours per Week column can be group into bins. b) In column Capital Gain, Capital Loss and Final Weight the data is left skewed.

```
In [8]: plt.figure(figsize=(10,10))
      sns.heatmap(data.corr(),annot = True)
```

Out[8]: <AxesSubplot:>



From the above graph we can observe that there is not very high linear correlation with the target column("Income") and final wweight column have zero correlation with the target value. so we can drop it.

AGE

bucketing the age columns into different bins as shown below


- a) 0-25: Young
- b) 25-50: Adult
- c) 50-100: Old

```
In [9]: data['Age'] = pd.cut(data['Age'], bins = [0, 25, 50, 100], labels = ['Young', 'Adult', 'Old'])
```

```
In [10]:
```

```
sns.countplot(x = 'Age', hue = 'Income', data = data)
```

```
Out[10]: <AxesSubplot:xlabel='Age', ylabel='count'>
```



From the above graph it can be observed that there is very less young people having Income grater than \$50K

Hours per Week

bucketing the age columns into different bins as shown below

a) 0-30: Low Hrs

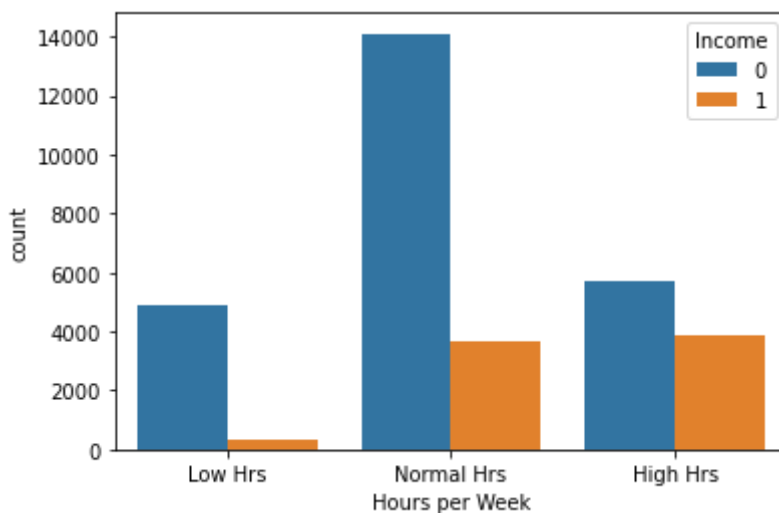
b) 30-40: Normal Hrs

c) 40-100: High Hrs

```
In [11]: data['Hours per Week'] = pd.cut(data['Hours per Week'], bins = [0, 30, 40, 100], lab
```

```
In [12]: sns.countplot(x = 'Hours per Week', hue = 'Income', data = data)
```

```
Out[12]: <AxesSubplot:xlabel='Hours per Week', ylabel='count'>
```



From the above graph it can be observe that as the working hours increase the no of people receiving more than 50k also increase.

Final Weight

As already discussed we can drop the column as it have zero correlation with the target column.

In [13]:

```
data.drop(['Final Weight'], axis = 1, inplace = True)
```

In [14]:

```
data
```

Out[14]:

	Age	Work Class	Education	Education Number	Marital Status	Occupation	Relationship	Race	Sex	Ca
0	Adult	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Adult	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Adult	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Old	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Adult	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	
...	
32556	Adult	Private	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	
32557	Adult	Private	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	
32558	Old	Private	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	
32559	Young	Private	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	
32560	Old	Self-emp-inc	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	1

32561 rows × 14 columns



Capital Gain and Capital loss


Using the above two column we can come up with new feature. Capital_diff= Capital Gain- Capital loss

In [15]:

```
data['Capital Diff'] = data['Capital Gain'] - data['Capital Loss']
data.drop(['Capital Gain'], axis = 1, inplace = True)
```

```
data.drop(['Capital Loss'], axis = 1, inplace = True)
data['Capital Diff'] = pd.cut(data['Capital Diff'], bins = [-5000, 5000, 100000], labels = ['low', 'high'])
sns.countplot(x = 'Capital Diff', hue = 'Income', data = data)
```

Out[15]: <AxesSubplot:xlabel='Capital Diff', ylabel='count'>




we can observe that for both the category low and high people with high Income (more than 50k)

Work Class

In [16]:

```
plt.figure(figsize=(20,10))
sns.countplot(x = 'Work Class', hue = 'Income', data = data)
```

Out[16]: <AxesSubplot:xlabel='Work Class', ylabel='count'>



From the above graph it can be observed that "without pay" and "Never-worked" column have very less records so it is safe to remove them. we can also observe a category "?", it is a error and have very low vaalue so we can remove it too.

In [17]:

```
data = data.drop(data[data['Work Class'] == '?'].index)
data = data.drop(data[data['Work Class'] == 'Without-pay'].index)
data = data.drop(data[data['Work Class'] == 'Never-worked'].index)
```


In [18]:

data

Out[18]:

	Age	Work Class	Education	Education Number	Marital Status	Occupation	Relationship	Race	Sex	Income
0	Adult	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	No
1	Adult	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	No
2	Adult	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	No
3	Old	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	No
4	Adult	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	No
...
32556	Adult	Private	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	No
32557	Adult	Private	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	No
32558	Old	Private	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	No
32559	Young	Private	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	No
32560	Old	Self-emp-inc	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	No

30704 rows × 13 columns




Education and Education Number

In [19]:

plt.figure(figsize=(20,10))
sns.countplot(x = 'Education', hue = 'Income', data = data)

Out[19]:

<AxesSubplot:xlabel='Education', ylabel='count'>



In [20]: `data['Education'].value_counts()`

Out[20]:

HS-grad	9959
Some-college	6772
Bachelors	5182
Masters	1675
Assoc-voc	1321
11th	1056
Assoc-acdm	1019
10th	831
7th-8th	572
Prof-school	558
9th	463
Doctorate	398
12th	393
5th-6th	303
1st-4th	156
Preschool	46

Name: Education, dtype: int64

Marital Status and Relationship

In [21]: `data['Relationship'].value_counts()`

Out[21]:

Husband	12700
Not-in-family	7865
Own-child	4520
Unmarried	3269
Wife	1432
Other-relative	918

Name: Relationship, dtype: int64

In [22]: `data['Marital Status'].value_counts()`

Out[22]:


Married-civ-spouse	14331
Never-married	9908
Divorced	4258
Separated	959
Widowed	839
Married-spouse-absent	388

Married-AF-spouse 21
 Name: Marital Status, dtype: int64

Occupation

```
In [23]: plt.figure(figsize=(20,10))
plt.xticks(rotation = 45)
sns.countplot(x = 'Occupation', hue = 'Income', data = data)
```

```
Out[23]: <AxesSubplot:xlabel='Occupation', ylabel='count'>
```




From the above graph it can be observe that there is no missing value and there is all unique categories so we can keep it as it is

Race

```
In [24]: plt.figure(figsize=(20,10))
sns.countplot(x = 'Race', hue = 'Income', data = data)
```

```
Out[24]: <AxesSubplot:xlabel='Race', ylabel='count'>
```




From the above graph it can be observed that the maximum information is about white people so we can combine other categories in one as others.

```
In [25]: data['Race'].replace([' Black', ' Asian-Pac-Islander', ' Amer-Indian-Eskimo', ' Othe
```

```
In [26]: plt.figure(figsize=(20,10))
sns.countplot(x = 'Race', hue = 'Income', data = data)
```


```
Out[26]: <AxesSubplot:xlabel='Race', ylabel='count'>
```



Sex

```
In [27]: plt.figure(figsize=(20,10))
sns.countplot(x = 'Sex', hue = 'Income', data = data)
```

```
Out[27]: <AxesSubplot:xlabel='Sex', ylabel='count'>
```



From the above it can be observe that there are more Male than compare to Female and there are more male reciving salary more than 50k

Country

In [28]: `data['Country'].value_counts()`

Out[28]:

United-States	27491
Mexico	610
?	556
Philippines	187
Germany	128
Puerto-Rico	109
Canada	107
India	100
El-Salvador	100
Cuba	92
England	86
Jamaica	80
South	71
China	68
Italy	68
Dominican-Republic	67
Vietnam	64
Guatemala	63
Japan	59
Poland	56
Columbia	56
Iran	42
Taiwan	42
Haiti	42
Portugal	34
Nicaragua	33
Peru	30
Greece	29
France	27
Ecuador	27
Ireland	24
Hong	19
Cambodia	18
Trinidad&Tobago	18
Thailand	17

```

Laos 17
Yugoslavia 16
Outlying-US(Guam-USVI-etc) 14
Hungary 13
Honduras 12
Scotland 11
Holand-Netherlands 1
Name: Country, dtype: int64

```

From the above it can be observed that there is some category with "?" that can be dropped and majority of people are from "US" so we can create two category as " US and OTHERS"

```

In [33]: dataset = data.drop(data[data['Country'] == '?'].index)
countries = np.array(data['Country'].unique())
countries = np.delete(countries, 0)
data['Country'].replace(countries, 'Other', inplace = True)

```

```


In [34]: plt.figure(figsize=(20,10))
sns.countplot(x = 'Country', hue = 'Income', data = dataset)

```

```

Out[34]: <AxesSubplot:xlabel='Country', ylabel='count'>

```



Splitting the datasets into features and target value

```

In [35]: data

```

```

Out[35]:

```

	Age	Work Class	Education	Education Number	Marital Status	Occupation	Relationship	Race	Sex	Income
0	Adult	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	No
1	Adult	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	No
2	Adult	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	No

	Age	Work Class	Education	Education Number	Marital Status	Occupation	Relationship	Race	Sex	
3	Old	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Other	Male	No
4	Adult	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Other	Female	No
...
32556	Adult	Private	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	No
32557	Adult	Private	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	No
32558	Old	Private	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	No
32559	Young	Private	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	
32560	Old	Self-emp-inc	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	No

30704 rows × 13 columns



```
In [37]: y = data['Income']
x= data.drop(['Income'], axis = 1)
x = pd.get_dummies(x)
print("Total features: {}".format(x.shape[1]))
```

Total features: 64


```
In [38]: x
```

Out[38]:

	Education Number	Age_Young	Age_Adult	Age_Old	Work Class_Federal-gov	Work Class_Local-gov	Work Class_Private	Work Class_Self-emp-inc	Work Class_Self-emp-not-inc	Work Class_State-gov
0	13	0	1	0	0	0	0	0	0	1
1	13	0	1	0	0	0	0	0	1	0
2	9	0	1	0	0	0	1	0	0	0
3	7	0	0	1	0	0	1	0	0	0
4	13	0	1	0	0	0	1	0	0	0
...
32556	12	0	1	0	0	0	1	0	0	0

	Education Number	Age_Young	Age_Adult	Age_Old	Work Class_ Federal- gov	Work Class_ Local- gov	Work Class_ Private	Work Class_ Self- emp- inc	Work Class_ Self- emp- not- inc	Work Class_ State- gov
32557	9	0	1	0	0	0	1	0	0	0
32558	9	0	0	1	0	0	1	0	0	0
32559	9	1	0	0	0	0	1	0	0	0
32560	9	0	0	1	0	0	0	1	0	0

30704 rows × 64 columns



In [39]:

y

Out[39]:

```
0      0
1      0
2      0
3      0
4      0
..
32556   0
32557   1
32558   0
32559   0
32560   1
Name: Income, Length: 30704, dtype: int32
```

In [41]:

```
from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.30, random_s
```

Machine learning

Importing libraries

In [42]:

```
from sklearn.metrics import f1_score, accuracy_score
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
```

In [43]:

```
classifiers = [GaussianNB(),
                SVC(kernel = 'rbf', probability = True),
                DecisionTreeClassifier(random_state = 0),
                RandomForestClassifier(n_estimators = 100, random_state = 0),
                GradientBoostingClassifier(random_state = 0)]
classifier_names = ["Gaussian Naive Bayes",
                    "Support Vector Classifier",
                    "Decision Tree Classifier",
                    "Random Forest Classifier",
                    "Gradient Boosting Classifier"]

accuracies = []
```

```
In [45]: for i in range(len(classifiers)):
          classifier = classifiers[i]
          classifier.fit(x_train, y_train)
          y_pred = classifier.predict(x_test)
          print("{}:".format(classifier_names[i]))
          print("F1 score: {:.2f}".format(f1_score(y_test, y_pred)))
          accuracy = accuracy_score(y_test, y_pred)*100
          accuracies.append(accuracy)
```

Gaussian Naïve Bayes:

F1 score: 0.61

Support Vector Classifier:

F1 score: 0.63

Decision Tree Classifier:

F1 score: 0.61

Random Forest Classifier:

F1 score: 0.65

Gradient Boosting Classifier:

F1 score: 0.66

From the above result it can be observe that "Gradient Boosting Classifier: F1 score: 0.66" is performing best

In []: